



The Style of a Machine : Distinguishing Humans from LLMS through Linguistic Fingerprinting

Hyungbin Park , Kihoon Lee, Gwangsu Kim
Department of Physics, Research Institute for Materials and Energy Sciences, Jeonbuk National University, Jeonju 54896, Republic of Korea

Introduction

Key question: Can we still tell apart human and machine-generated text?

- Modern LLMS(e.g., GPT, Gemini) produce text nearly indistinguishable from human writing
- This study explores “Stylistic Differences” between human responses and those from major LLMS
- Focused on outputs from Claude, DeepSeek, Gemini, GPT and Human
- Aims to identify “Unique Stylistic Fingerprints” for each source
- Investigates whether subtle yet detectable patterns exit despite fluent coherent writing

Data & Methodology

- First, for the purpose of classifying responses as either LLM-generated or human-written, we collected real-world answers from human users on Quora and generated responses from four major LLMS: Claude, Gemini, DeepSeek and GPT
- Second, we used five distinct prompt modes to efficiently analyze data diversity and varied LLM response characteristics
- We explored various linguistic features (e.g., stylistic stats, POS tags, TF-IDF, SBERT, LDA) to analyze the responses.
- For classification, we trained a LightGBM model using only the LDA-based subtopic distributions as input features.

Prompt Typology

Statistical Analysis

Minimal Guidance

- Uses minimal, neutral instructions to capture the LLM's default response style

Objective Response

- Directs the LLM to prioritize factual, objective information, testing its ability

Experience Sharing

- Aims to elicit subjective, experience-based narratives and opinions, assessing the LLM's capacity for a personal

Analytical Response

- Encourages detailed, analytical answers pushing the LLM to explore topics comprehensively

TF-IDF, SBERT, Classification

Core Informative & Analytical Set


- Assesses versatile LLM responses by combining Minimal, Analytical, objective.
- Subjective experience prompts were excluded to first establish a standardized base-Line for core information processing capabilities.
- A total of 100 questions were curated across multiple topics, with 30 responses per question sourced from four LLMS and human contributors. While LLM responses were consistently sampled to provide 30 answers per question, the number of human responses varied, and some questions had fewer than 30 human answers available.

- Lastly, the input texts were cleaned by removing special characters, lowercased, tokenized using NLTK(Natural Language Tool Kit), and filtered to exclude English stopwords for data preprocessing & tokenization


Stylistic Fingerprinting : How They Speak

- Linguistic analysis(stylometry, POS) clearly distinguishes human text from LLMS. Human show vast stylistic diversity. LLMS exhibit more consistent, model-specific signatures.


Char Count



Lexical Density



Unique Word Count



Topical & Lexical Signatures : What They Speak about


Term Frequency - Inverse Document Frequency

- LLMS generate responses using a common, general-purpose lexicon. In contrast Human lexical choices indicate responses more grounded in specific contexts, personal perspectives, temporal awareness

Top-10	Human		GPT		Gemini		DeepSeek		Claude	
	Word	Score	Word	Score	Word	Score	Word	Score	Word	Score
1	One	0.2570	Like	0.3072	Like	0.3796	Like	0.3088	Like	0.2530
2	Time	0.2424	One	0.2373	Think	0.2711	One	0.2925	People	0.2247
3	Would	0.2235	Time	0.1840	Thing	0.2285	People	0.2696	I've	0.2189
4	People	0.1879	People	0.1837	One	0.1966	Even	0.2426	Think	0.2042
5	Like	0.1807	Think	0.1768	People	0.1805	Way	0.1815	Something	0.1804
6	Get	0.1582	Might	0.1735	Even	0.1754	Thing	0.1621	One	0.1750
7	Day	0.1550	Thing	0.1595	Okay	0.1751	Time	0.1506	Actually	0.1681
8	Year	0.1521	I've	0.1563	Really	0.1610	Make	0.1402	Thing	0.1529
9	Life	0.1457	Even	0.1517	Know	0.1607	Something	0.1389	Time	0.1517
10	dont	0.1381	Life	0.1473	something	0.1361	dont	0.1331	someone	0.1344

Sentence BERT

- SBERT analysis shows LLMS are highly similar and consistent in their semantic meaning.
- Human responses are distinctly different from LLMS and far more internally diverse



Classification

Hierarchical LDA-Based Subtopic Modeling

- Topic Modeling (LDA)
 - Construct Dictionary and Bag-of-Words from tokenized data
 - The number of topics was set to five(four LLMS and one human)
 - Represent each document as a distribution over main topics
- Subtopic Modeling
 - Select top 20 representative documents for each main topic
 - Identify optimal number of subtopics within each document group
 - Apply subtopic model to all documents to compute subtopic probability distributions

Doc_id	0-1	0-2	1-1	1-2	1-3	1-4	...	LLM_Name
1	0.7245	0.2755	0.1853	0.2493	0.2679	0.2975	...	GPT-4
2	0.9074	0.0926	0.2296	0.2947	0.1642	0.3115	...	Claude
3	0.8692	0.1308	0.2779	0.2186	0.2379	0.2656	...	Claude
4	0.8739	0.1261	0.1452	0.2068	0.3081	0.3399	...	Deepseek
...								
n-1	0.8486	0.1514	0.2454	0.2154	0.3099	0.2292	...	Human
n	0.8790	0.1210	0.1651	0.2923	0.2491	0.2935	...	Human

* This matrix is used as the input feature set for classification

LightGBM Based Classification

- Input X : Subtopic probability vector per document
- Target Y : LLM name associated with each document
- Classifier : LightGBM multi-class model (Train 60% : Test 40%)

	precision	recall	f1-score	support
claude	0.614907	0.66	0.636656	1200
deepseek	0.679865	0.6725	0.676163	1200
gemini	0.75136	0.805833	0.777644	1200
gpt	0.805461	0.786667	0.795953	1200
human	0.86	0.758377	0.805998	1134
accuracy	0.736434			
Macro avg	0.742319	0.736675	0.738483	5934
Weighted avg	0.74101	0.736434	0.737732	5934

LightGBM Confusion Matrix

- The LightGBM model had an overall accuracy of 73.6%, classifying human responses most accurately, while Claude and DeepSeek were often confused. In particular, Gemini and GPT had high f1-scores, showing that they reflect the differences in responses between LLMS.

Conclusion & Future Work

- This study demonstrates that distinctive linguistic fingerprints exit, enabling reliable classification between human and LLM-generated text, as well as among different LLMS. These differences stem from core features like semantic diversity and narrative style.
- The classifier developed here can help combat Model Collapse the degradation of model trained on AI generated content. By assigning an “AI-generated probability score” to training data, it can assist in filtering such content, preserving data diversity and ensuring more sustainable LLM development.

Acknowledgement

This research was supported by Global – Learning & Academic research institution for Master’s-PhD students, and Postdocs(LAMP) Program of the National Research Foundation of Korea(NRF) grant funded by the Ministry of Education (No. RS-2024-00443714).